

DM2: Probabilistic Graphical Models

4 Implementation - Gaussian Mixture

```

In [1]: %matplotlib inline
        from __future__ import division

        import pandas as pd
        import numpy as np
        from scipy.stats import norm, multivariate_normal
        import math
        import random

        import matplotlib.pyplot as plt
        from matplotlib.patches import Ellipse
        import matplotlib.cm as cm

        plt.style.use('ggplot')
        plt.rcParams['figure.figsize'] = 12,12

        em_train = pd.read_csv('EMGaussian.data',header=None,sep=' ')
        em_test = pd.read_csv('EMGaussian.test',header=None,sep=' ')

        # Inspired from http://stackoverflow.com/questions/12301071/multidimensional-confidence-intervals
        def plot_cov_ellipse(cov, pos, nstd=1.645, ax=None, fill=False, **kwargs):
            """
            Plots an `nstd` sigma error ellipse based on the specified covariance
            matrix (`cov`). Additional keyword arguments are passed on to the
            ellipse patch artist.

            Parameters
            -----
            cov : The 2x2 covariance matrix to base the ellipse on
            pos : The location of the center of the ellipse. Expects a 2-element
                  sequence of [x0, y0].
            nstd : The radius of the ellipse in numbers of standard deviations.
                  Defaults to 1.645 standard deviations (for 90% of the mass
                  of the Gaussian distribution)
            ax : The axis that the ellipse will be plotted on. Defaults to the
                  current axis.
            Additional keyword arguments are pass on to the ellipse patch.

            Returns
            -----
            A matplotlib ellipse artist
            """
            def eigsorted(cov):
                vals, vecs = np.linalg.eigh(cov)
                order = vals.argsort()[::-1]
                return vals[order], vecs[:,order]

            if ax is None:
                ax = plt.gca()

            vals, vecs = eigsorted(cov)

```

```

theta = np.degrees(np.arctan2(*vecs[:,0][::-1]))

# Width and height are "full" widths, not radius
width, height = 2 * nstd * np.sqrt(vals)
ellip = Ellipse(xy=pos, width=width, height=height, angle=theta, fill=False, **kwargs)

ax.add_artist(ellip)
return ellip

```

(a). Implementation of kmeans

```

In [2]: def closest_center(point, centers):
        centers = np.asarray(centers)
        dist_2 = np.sum((centers - point)**2, axis=1)
        return np.argmin(dist_2), np.min(dist_2)

def kmeans(df, k):
    n = len(df)
    indexes = random.sample(range(n), k)
    new_centers = np.array([df.iloc[i].values for i in indexes])
    centers = np.array([])
    distorsion = []

    while not np.array_equal(centers, new_centers):
        centers = new_centers

        # Assign each point to the closest center
        closest = np.array([list(closest_center(x.values, centers)) for _, x in df.iterrows()])
        labels = closest[:,0]
        distorsions = closest[:,1]
        distorsion.append(distorsions.sum())

        # Compute new centers
        new_centers = np.array([df[labels == i].mean(axis=0) for i in range(k)])
    return new_centers, labels, distorsion


```


K-means is applied several times with randoms initialisation, the clustering with the center is displayed along with the distorsion at each iteration


```
In [3]: k = 4
color_map = cm.get_cmap('coolwarm')
colors = color_map([x/float(k) for x in range(k)])
for i in range(0,3):
    centers, labels, distorsion = kmeans(em_train, k)


    ax = em_train[labels == 0].plot(kind='scatter', x=0, y=1, color=colors[0], s=80)
    for j in range(1,k):
        em_train[labels == j].plot(kind='scatter', x=0, y=1, ax=ax, color=colors[j], s=80)
    plt.plot(centers[:,0], centers[:,1], 'o', ms=13, color='red')
    plt.title('K-means clustering with k={}'.format(k))
    plt.show()
    plt.plot(distorsion)
    plt.title('Distorsion for each iteration of K-means')
    plt.show()
```


```
/Users/Adrien/anaconda/lib/python2.7/site-packages/matplotlib/collecti  
ons.py:571: FutureWarning: elementwise comparison failed; returning sc  
alar instead, but in the future will perform elementwise comparison  
if self._edgecolors == str('face'):
```














The clustering is efficient with four cluster and doesn't seem to suffer from random initialization

(c). Gaussian mixture

```

In [4]: # The covariance matrices are proportional the identity matrix
def em_identity(df, K, iter=100):
    array = df.values
    n = len(df)
    p = np.empty(K); p.fill(1/n)
    sigma = [1] * 4

    # Initialize centers with kmeans
    mu, _, _ = kmeans(df, K)
    mu = list(mu)
    res = np.empty((K, n))

    for j in xrange(iter):
        # E step
        for k in range(K):
            res[k,:] = p[k] * multivariate_normal.pdf(array,mu[k],sigma
a[k])
        res = res / res.sum(axis=0)

        # M step
        p = res.mean(axis=1)
        # The responsibilities are normalized to avoid dividing by the
ir sum for the next computations
        res = (res.T / res.sum(axis=1)).T
        mu = np.dot(res,array)
        for k in range(K):
            sigma[k] = res[k,:].dot(np.linalg.norm(array - mu[k],axi
s=1))
    return p, mu, sigma

def predict(df, p, mu, sigma):
    array = df.values
    n = len(df)
    k = len(mu)
    proba = np.zeros((k,n))
    for i in xrange(k):
        proba[i,:] = p[i] * multivariate_normal.pdf(array, mu[i], sigma
a[i])
    return np.argmax(proba, axis=0)

```


```

In [5]: k = 4
data = em_train
color_map = cm.get_cmap('coolwarm')
colors = color_map([x/float(k) for x in range(k)])

p1, mu1, sigma1 = em_identity(em_train, k)
sigma1 = [np.diag(np.repeat(s,2)) for s in sigma1]
labels = predict(em_train, p1, mu1, sigma1)

ax = data.plot(kind='scatter', x=0, y=1, color=colors[0], s=80)
plot_cov_ellipse(sigma1[0], mu1[0], color='red', fill=False)
for j in range(1,k):
    data[labels == j].plot(kind='scatter', x=0, y=1, ax=ax, color=colors[j], s=80)
    plot_cov_ellipse(sigma1[j], mu1[j], color='red', fill=False)
plt.plot(mu1[:,0], mu1[:,1], 'o', ms=13, color='red')
plt.title('EM clustering circle covariances with K={}'.format(k))
# plt.gca().set_aspect('equal', adjustable='box') # used to ensure axes have same scale
plt.show()

```



Here cluster are circles (you can uncomment the last line to have equal axis on the graph), we can observe that it doesn't seem to be right model to be fit to our data

```
In [6]: def em(df, K, iter=100):
    array = df.values
    n = len(df)
    p = np.empty(K); p.fill(1/n)
    sigma = [df.cov().values] * K

    # Initialize centers with kmeans
    mu, _, _ = kmeans(df, K)
    mu = list(mu)
    res = np.empty((K, n))

    for j in xrange(iter):
        # E step
        for k in range(K):
            res[k,:] = p[k] * multivariate_normal.pdf(array, mu[k], sigma
a[k])
        res = res / res.sum(axis=0)

        # M step
        p = res.mean(axis=1)
        # The responsibilities are normalized to avoid dividing by the
ir sum for the next computations
        res = (res.T / res.sum(axis=1)).T
        mu = np.dot(res, array)
        for k in range(K):
            sigma[k] = (res[k,:] * (array - mu[k]).T).dot(array - m
u[k])
    return p, mu, sigma
```


```

In [7]: k = 4
data = em_train
color_map = cm.get_cmap('coolwarm')
colors = color_map([x/float(k) for x in range(k)])

p2, mu2, sigma2 = em(em_train, k)
labels = predict(em_train, p2, mu2, sigma2)

ax = data.plot(kind='scatter', x=0, y=1, color=colors[0], s=80)
plot_cov_ellipse(sigma2[0], mu2[0], color='red', fill=False)
for j in range(1,k):
    data[labels == j].plot(kind='scatter', x=0, y=1, ax=ax, color=colors[j], s=80)
    plot_cov_ellipse(sigma2[j], mu2[j], color='red', fill=False)
plt.plot(mu2[:,0], mu2[:,1], 'o', ms=13, color='red')
plt.title('EM clustering general covariance matrices with K={}'.format(k))
plt.show()

```



This time with general covariance matrix the model fit better the different shapes of our data

(d). Loglikelihood of the model on train data and test on test data

As expected from the plot of the data and the shape of the clusters the second model with general covariance matrices is performing better for clustering.

```
In [8]: def loglikelihood(df, mu, sigma, p):  
        K = len(p)  
        ll = np.zeros((len(df), K))  
        array = df.values  
        for k in xrange(0,K):  
            ll[:,k] = multivariate_normal.pdf(data, mu[k,:], sigma[k])  
        ll = np.dot(ll,p)  
        return np.log(ll).sum()
```

```
In [11]: print "Log-likelihood with model 1 (circle covariance matrices) on tra  
in data: {}".format(loglikelihood(em_train, mu1, sigma1,p1))  
print "Log-likelihood with model 2 (general covariance matrices) on tr  
ain data: {}\n".format(loglikelihood(em_train, mu2, sigma2,p2))  
print "Log-likelihood with model 1 on test data: {}".format(loglikelih  
ood(em_test, mu1, sigma1,p1))  
print "Log-likelihood with model 2 on test data: {}".format(loglikelih  
ood(em_test, mu2, sigma2,p2))
```

```
Log-likelihood with model 1 (circle covariance matrices) on train dat  
a: -2703.77401346  
Log-likelihood with model 2 (general covariance matrices) on train dat  
a: -2327.71567492
```

```
Log-likelihood with model 1 on test data: -2703.77401346  
Log-likelihood with model 2 on test data: -2327.71567492
```

As expected the Log-likelihood is higher with the second model